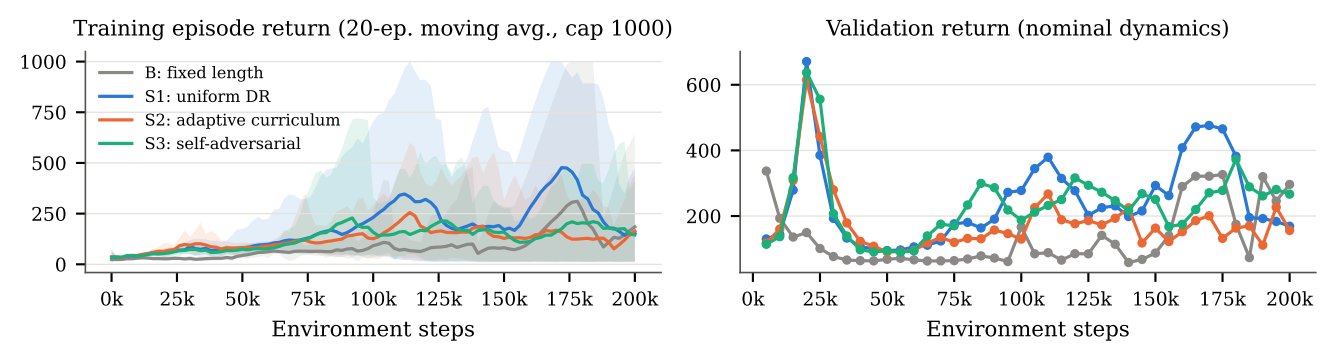
Training episode return (20-ep. moving avg., cap 1000)
B: fixed length
S1: uniform DR
S2: adaptive curriculum
S3: self-adversarial
1000
750
500
250
0
0k
25k
50k
75k
100k
125k
150k
175k
200k
Environment steps
Validation return (nominal dynamics)
600
400
200
0k
25k
50k
75k
100k
125k
150k
175k
200k
Environment steps